

Analysis of Key Statistics in Asset Returns

```
[27]: import os
import time
import datetime
import numpy as np
import pandas as pd
import scipy.stats as scs
from pylab import plt, mpl

plt.style.use('seaborn-v0_8')
mpl.rcParams['savefig.dpi'] = 300
mpl.rcParams['font.family'] = 'serif'
pd.set_option('mode.chained_assignment', None)
pd.set_option('display.float_format', '{:.4f}'.format)
np.set_printoptions(suppress=True, precision=4)
os.environ['PYTHONHASHSEED'] = '0'
```

An In-depth Analysis

In the realm of financial theories, the normality assumption plays a pivotal role in dealing with uncertain returns of financial instruments. In this section, we analyze four historical financial time series using real-world data.

Data

We will analyze four historical financial time series, two for technology stocks and two for exchange traded funds (ETFs).

```
[28]: dataFrame = pd.read_csv('data_.csv',
                             index_col=0, parse_dates=True).dropna()
columns = ['SPY', 'GLD', 'AAPL.O', 'MSFT.O']
dFrame = dataFrame[columns].dropna()
dFrame.info() # display dataset information
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 2516 entries, 2010-01-04 to 2019-12-31
Data columns (total 4 columns):
#   Column  Non-Null Count  Dtype
---  -
0    SPY      2516 non-null    float64
```

```

1   GLD      2516 non-null   float64
2   AAPL.O   2516 non-null   float64
3   MSFT.O   2516 non-null   float64
dtypes: float64(4)
memory usage: 98.3 KB

```

```
[29]: dFrame.head() # display the first five rows
```

```

[29]:          SPY      GLD  AAPL.O  MSFT.O
Date
2010-01-04  113.3300  109.8000  30.5728  30.9500
2010-01-05  113.6300  109.7000  30.6257  30.9600
2010-01-06  113.7100  111.5100  30.1385  30.7700
2010-01-07  114.1900  110.8200  30.0828  30.4520
2010-01-08  114.5700  111.3700  30.2828  30.6600

```

Normalize the financial instruments' prices

Figure 1-1. shows the normalized prices of the financial assets

```

[30]: (dFrame / dFrame.iloc[0] * 100).plot(figsize=(10, 6))
plt.figtext(0.5, 0.0001, 'Fig 1-1. Normalized prices of the financial assets',
           style='italic', ha='center')
plt.show()

```



Fig 1-1. Normalized prices of the financial assets

Log returns of the financial instruments as histograms

```
[31]: log_returns = np.log(dFrame / dFrame.shift(1))
log_returns.dropna(inplace=True)
log_returns.head()
```

```
[31]:          SPY      GLD  AAPL.O  MSFT.O
Date
2010-01-05  0.0026 -0.0009  0.0017  0.0003
2010-01-06  0.0007  0.0164 -0.0160 -0.0062
2010-01-07  0.0042 -0.0062 -0.0019 -0.0104
2010-01-08  0.0033  0.0050  0.0066  0.0068
2010-01-11  0.0014  0.0132 -0.0089 -0.0128
```

Figure 1-2. shows the log returns of the financial instruments as histograms

```
[32]: log_returns.hist(bins=50, figsize=(10, 8));
plt.figtext(0.5, 0.0001, 'Fig 1-2. Histograms of log returns for financial_
↳instruments', style='italic', ha='center')
plt.show()
```

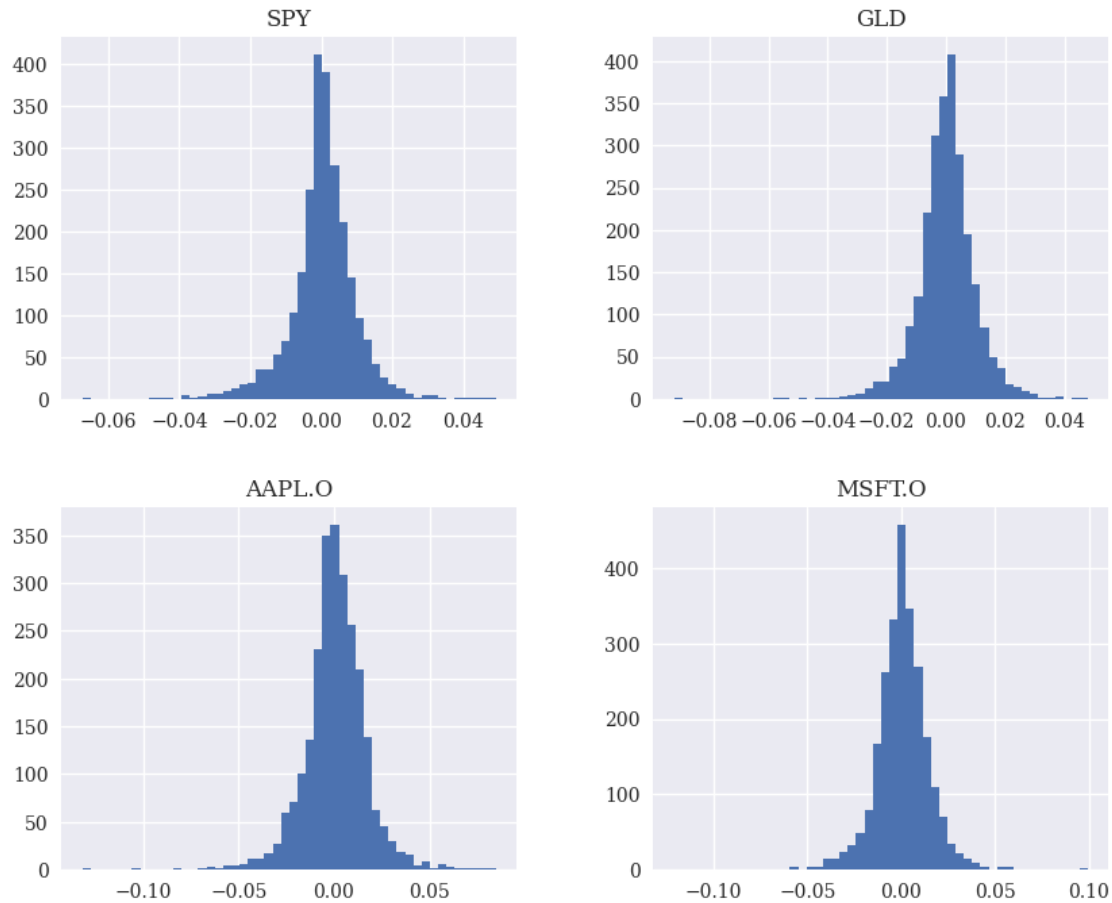


Fig 1-2. Histograms of log returns for financial instruments

Examining key statistics

We will use the function `print_statistics()`, which generates a better (human) readable key statistics output.

```
[33]: def print_statistics(array):
        '''Prints selected statistics.
        Parameters
        =====
        array: ndarray
               object to generate statistics on
        '''
        sta = scs.describe(array)
        print('%14s %15s' % ('statistic', 'value'))
        print(30 * '-')
        print('%14s %15.5f' % ('size', sta[0]))
```

```

print('%14s %15.5f' % ('min', sta[1][0]))
print('%14s %15.5f' % ('max', sta[1][1]))
print('%14s %15.5f' % ('mean', sta[2]))
print('%14s %15.5f' % ('std', np.sqrt(sta[3])))
print('%14s %15.5f' % ('skew', sta[4]))
print('%14s %15.5f' % ('kurtosis', sta[5]))

```

```

[34]: for col in columns:
    print('\nResults for column {}'.format(col))
    print(30 * '-')
    log_data = np.array(log_returns[col].dropna())
    print_statistics(log_data)

```

Results for column SPY

statistic	value
size	2515.00000
min	-0.06734
max	0.04929
mean	0.00042
std	0.00930
skew	-0.50079
kurtosis	4.45059

Results for column GLD

statistic	value
size	2515.00000
min	-0.09191
max	0.04795
mean	0.00010
std	0.00978
skew	-0.58102
kurtosis	5.89970

Results for column AAPL.O

statistic	value
size	2515.00000
min	-0.13187
max	0.08502
mean	0.00090
std	0.01625

skew	-0.33706
kurtosis	4.80934

Results for column MSFT.O

statistic	value
size	2515.00000
min	-0.12103
max	0.09941
mean	0.00065
std	0.01433
skew	-0.10594
kurtosis	6.41239

Analysis

Kurtosis, the fourth moment, characterizes the shape of the distribution's tails. It measures the peakedness or flatness of the distribution. In all four of our data sets, the high positive kurtosis results in a fatter tail on the right side, indicating a more pronounced peak compared to a normal distribution.

Normality Test

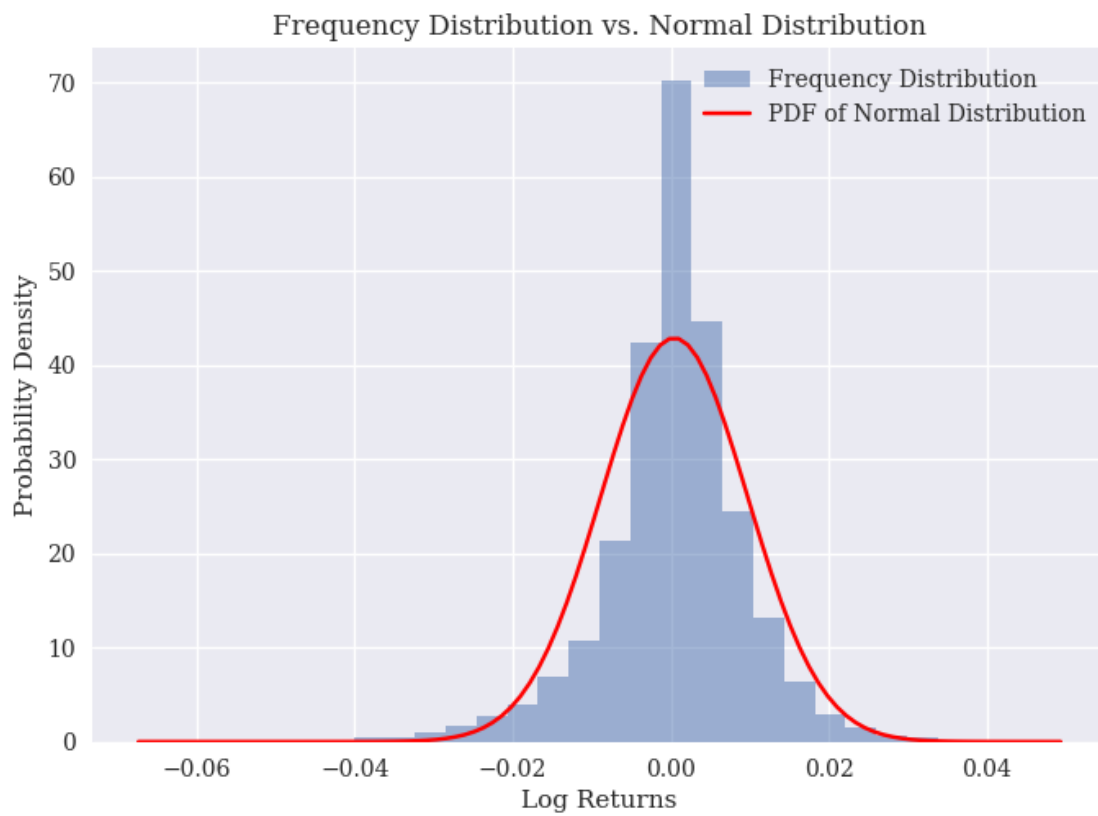
a. Graphical Normality Test

1. Compare the histograms and the PDF for normal distribution

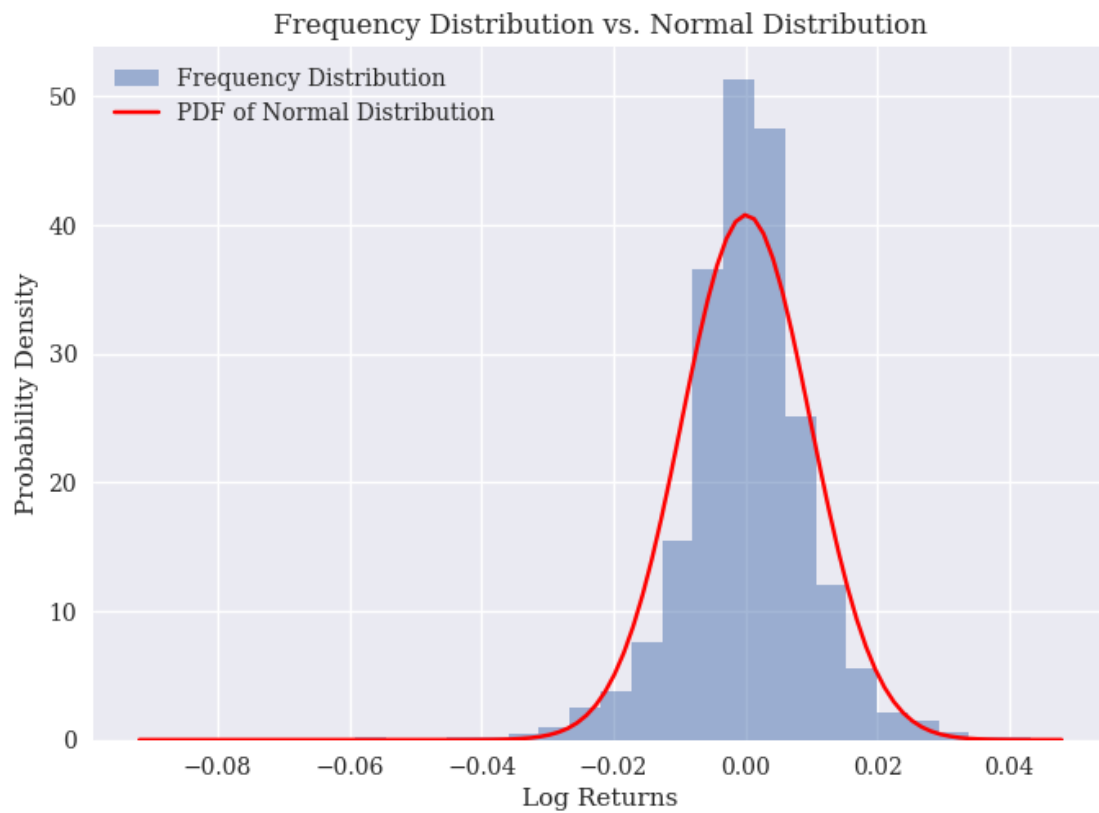
```
[35]: from scipy.stats import norm
for col in columns:
    print('\nResults for column {}'.format(col))
    print(30 * '-')
    log_data = np.array(log_returns[col].dropna())

    # Plotting histogram and PDF for the assets
    plt.hist(log_data, bins=30, density=True, alpha=0.5, label='Frequency_
↪Distribution')
    mu, sigma = log_data.mean(), log_data.std()
    x = np.linspace(log_data.min(), log_data.max(), 100)
    plt.plot(x, norm.pdf(x, mu, sigma), 'r-', label='PDF of Normal_
↪Distribution')
    plt.xlabel('Log Returns')
    plt.ylabel('Probability Density')
    plt.title('Frequency Distribution vs. Normal Distribution')
    plt.legend()
    plt.show()
```


Results for column SPY



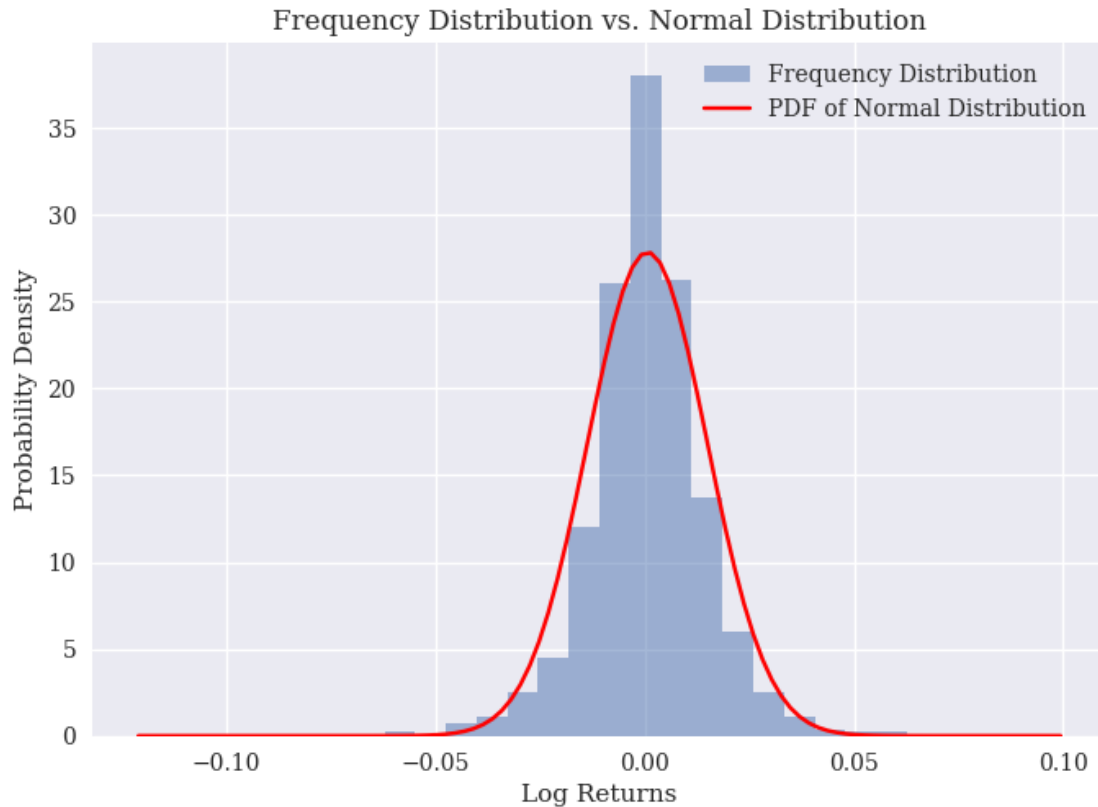
Results for column GLD



Results for column AAPL.0



Results for column MSFT.0



[]:

Analysis


The histograms above represent the frequency distribution of the log returns, while the PDF represents the expected distribution if the log returns followed a normal distribution. By visually comparing the histograms and the PDF, we can assess if there is a good fit between the log returns and the normal distribution. In our case, we observe a slightly left-skewed distribution, indicating a longer left tail compared to a normal distribution.

2. Quantile-quantile plots


```
[36]: import statsmodels.api as sm
for col in columns:
    print('\nResults for column {}'.format(col))
    print(30 * '-')
    log_data = np.array(log_returns[col].dropna())

    #Generate the Q-Q plot
    sm.qqplot(log_data, line='s')
    plt.show()
```


Results for column SPY




Results for column GLD



Results for column AAPL.O



Results for column MSFT.0



Analysis

The Q-Q plots displayed above depict the log returns, with the red line representing a perfect fit to a normal distribution. By visually examining the plot, we can assess the goodness of fit of the log returns to a normal distribution. It is evident that the points on the plot deviate from the reference line, indicating non-normality. On both the left and right sides, there are numerous values that lie significantly below and above the line, respectively. This observation provides evidence of a skewed or fat-tailed distribution, signifying a departure from normality.

b. Analytical Normality Test

However enticing the graphical approaches may appear, they are generally unable to substitute more rigorous testing procedures. In the following example, the function `normality_tests()` integrates three distinct statistical tests:

Skewness test (`skewtest()`): Tests whether the skew of the log returns is “normal” (i.e., has a value close enough to zero).

Kurtosis test (`kurtosistest()`): Similarly, this tests whether the kurtosis of the log returns is “normal” (again, close enough to zero).

Normality test (`normaltest()`): Combines the other two test approaches to test for normality.

```
[37]: def normality_tests(arr):
      ''' Tests for normality distribution.
```

```

Parameters
=====
array: ndarray
    object to generate statistics on
'''

print('Skew of data set %14.5f' % scs.skew(arr))
print('Skew test p-value %14.5f' % scs.skewtest(arr)[1])
print('Kurt of data set %14.5f' % scs.kurtosis(arr))
print('Kurt test p-value %14.5f' % scs.kurtosistest(arr)[1])
print('Norm test p-value %14.5f' % scs.normaltest(arr)[1])

for col in columns:
    print('\nResults for column {}'.format(col))
    print(30 * '-')
    log_data = np.array(log_returns[col].dropna())
    normality_tests(log_data)

```

Results for column SPY

```

-----
Skew of data set      -0.50079
Skew test p-value     0.00000
Kurt of data set      4.45059
Kurt test p-value     0.00000
Norm test p-value     0.00000

```

Results for column GLD

```

-----
Skew of data set      -0.58102
Skew test p-value     0.00000
Kurt of data set      5.89970
Kurt test p-value     0.00000
Norm test p-value     0.00000

```

Results for column AAPL.O

```

-----
Skew of data set      -0.33706
Skew test p-value     0.00000
Kurt of data set      4.80934
Kurt test p-value     0.00000
Norm test p-value     0.00000

```

Results for column MSFT.O

```

-----
Skew of data set      -0.10594
Skew test p-value     0.03005
Kurt of data set      6.41239

```

Kurt test p-value	0.00000
Norm test p-value	0.00000

Analysis

The log returns data with a skewness of $-(0.)$ and a kurtosis of $3+$ exhibits interesting characteristics.

The negative skewness indicates that the log returns' distribution is slightly left-skewed, meaning that the tail on the left side is more spread out than the tail on the right side, an indication there may be more extreme negative returns in the data compared to extreme positive returns.

The kurtosis value of $3+$ indicates heavy tails and a higher peak compared to a normal distribution, suggesting a higher likelihood of observing extreme returns.

The p-values obtained from the different tests are all zero, indicating a complete rejection of the test hypothesis, which suggests that the log returns do not follow a normal distribution. This finding challenges the validity of the normal assumption for stock market returns and other asset classes, such as those embodied in the geometric Brownian motion model. Consequently, it becomes apparent that the utilization of richer models, capable of generating fat tails, such as jump diffusion models or models with stochastic volatility, may be necessary in order to accurately capture the underlying dynamics.